

# **Introduction to Information Security**

**14-741/18-631 Fall 2021**

**Unit 4: Lecture 2:**

**TCP/IP Vulnerabilities**

**Limin Jia**

**liminjia@andrew**

# This lecture's agenda

## ■ Outline

- ▼ Brief networking review
- ▼ TCP attacks
- ▼ Routing attacks
- ▼ Vulnerable applications
- ▼ Defenses

## ■ Objective

- ▼ Examine some of the key vulnerabilities in the widely used TCP/IP and related protocols

# Origins of the Internet

- **U.S. military**

- ▼ (D)ARPAnet
- ▼ Most research done in the mid-60s through late 70s

- **Objective**

- ▼ Have a network capable of withstand massive failures
- ▼ e.g., nuclear war

- **TCP/IP designed with reliability and fault-tolerance in mind**

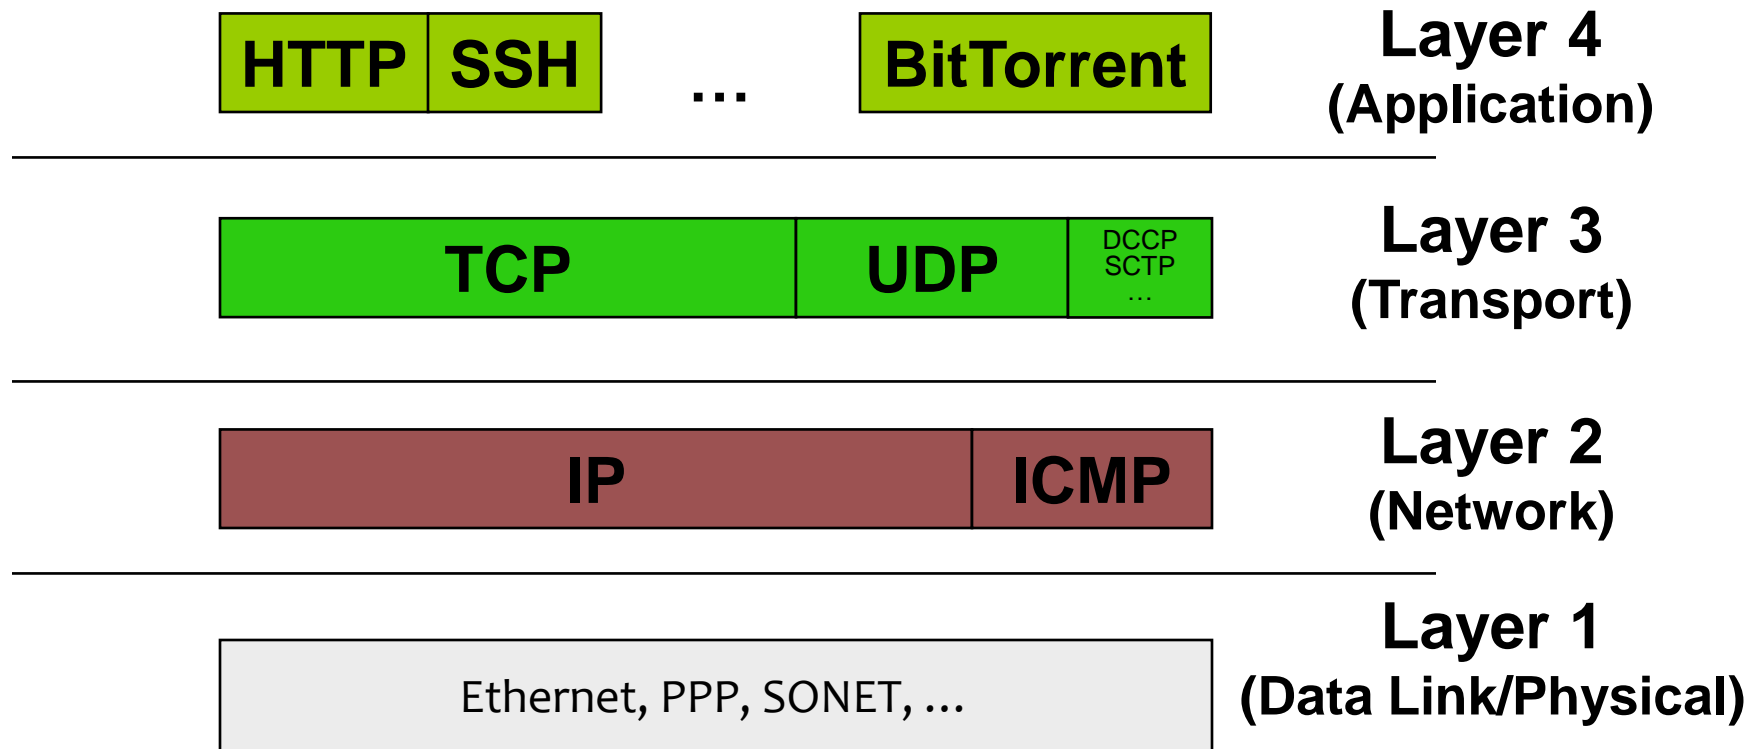
- **Does this automatically provide security?**

- **TCP/IP protocol suite overwhelmingly used on the Internet**
- **Layered architecture**
  - ▼ Different from 7-layer OSI model
  - ▼ Essentially no session or presentation layers
  - ▼ Physical and data link layer basically merged

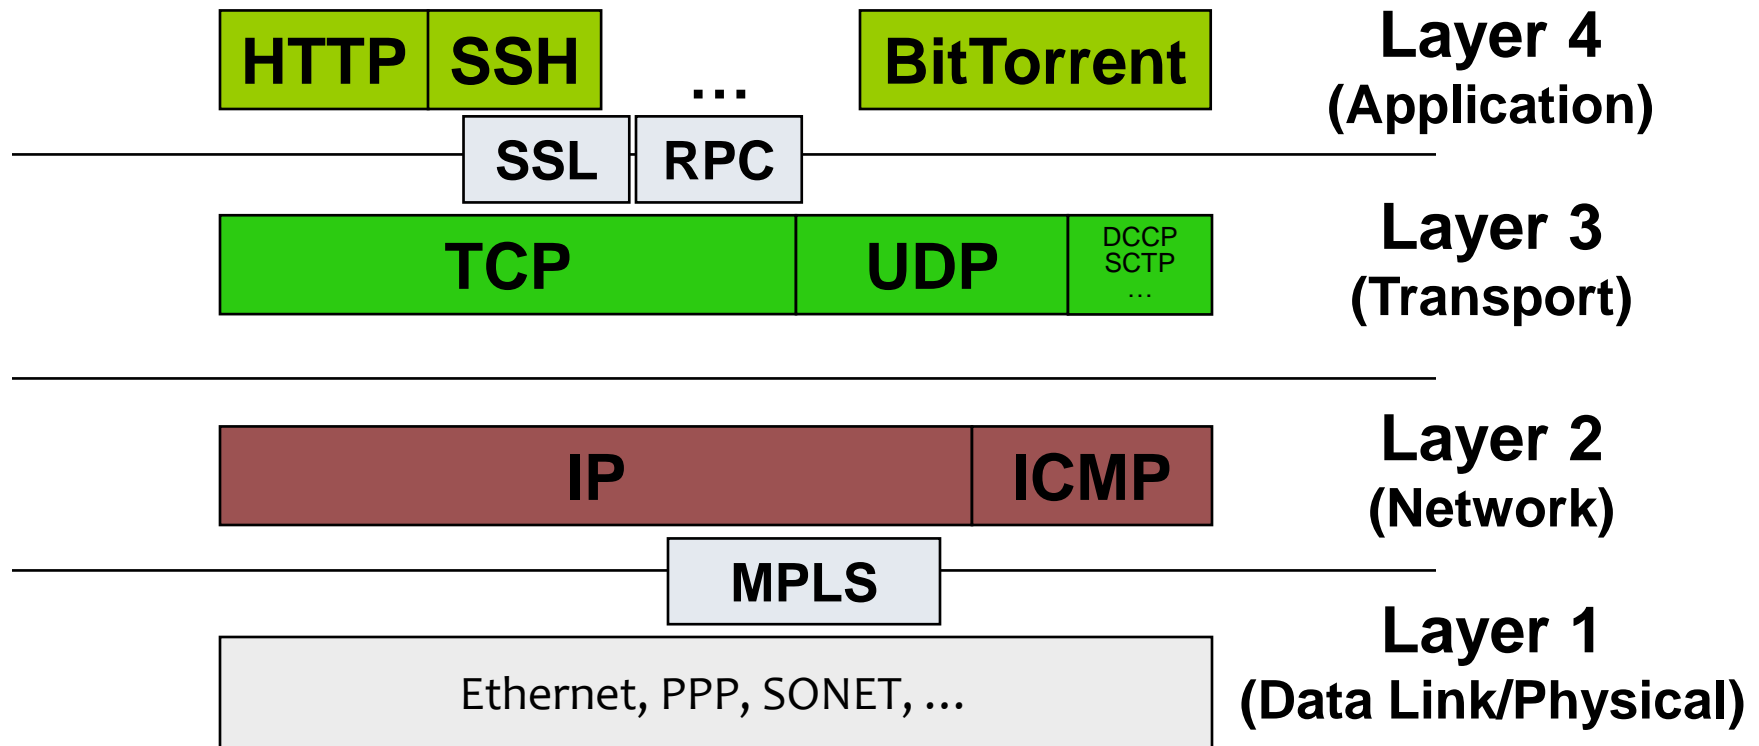
# Layered architecture

- Each layer provides service to layer above and relies on services from layer below
  - ▼ Why?
- Different functions at each layer
  - ▼ Data link: physical (local) connectivity
  - ▼ Network: end-to-end reachability
    - ▼ Relies on physical connectivity
  - ▼ Transport: reliability, ordering, ...
    - ▼ Relies on end-to-end reachability
- In general, layer **n** only talks with layers **n-1** and **n+1**
  - ▼ Exceptions (aberrations?): HTTP over IP, QUIC,...

# TCP/IP protocol stack




# TCP/IP protocol stack



# Data link layer: Local connectivity


## ■ Example: 802.11b





# Network layer: Global connectivity

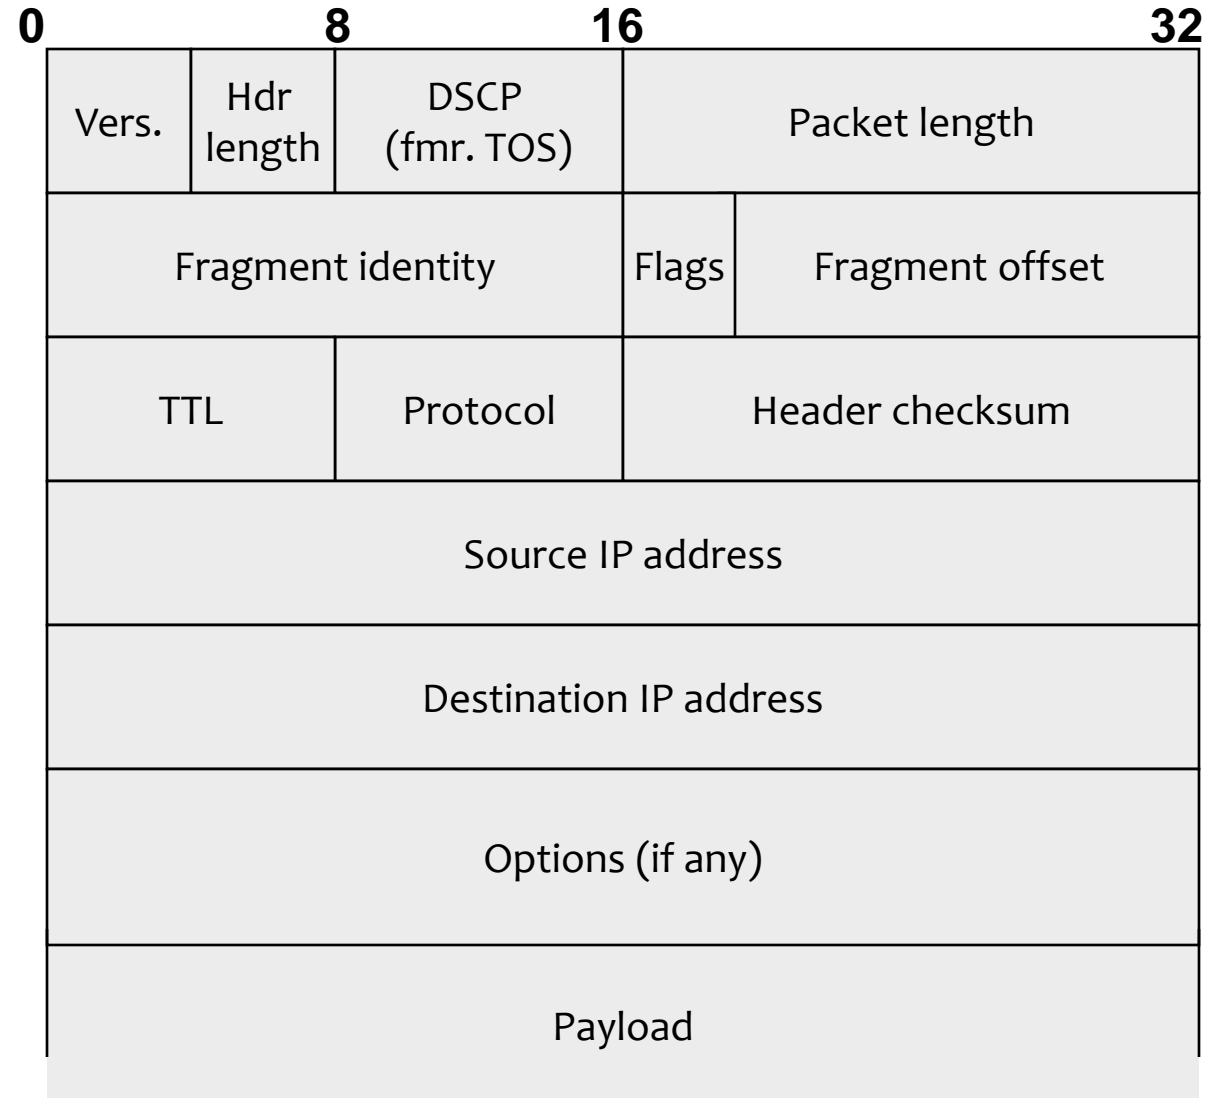
- **Routing:** How does the packet get from A to B?
- **Fragmentation**
  - ▼ Making sure each local medium (e.g., 802.11b, Ethernet, SONET) can accommodate all packets
  - ▼ Find smallest MTU and break up application layer packets accordingly



# More on routing

- **Forwarding tables at each router populated by routing protocols**
  - ▼ RIP, OSPF, BGP, ...
- **Originally (a long time ago), manually updated**
  - ▼ Humans still (too much?) in the loop
  - ▼ BGP misconfigurations are a serious problem
- **Routing protocols update tables based on cost of each route**
  - ▼ Exchange tables with neighbors, or everyone (depends on protocol)
  - ▼ Use neighbor leading to the shortest path

# IP packet format



# Transport layer: Reliability

## ■ Connection-oriented service

- ▼ Network layer only provides connectionless service

## ■ Provides resilience to

- ▼ Data corruption
  - ▼ Payload checksum
- ▼ Data loss
  - ▼ Sequence numbers, timeouts + retransmits
- ▼ Out-of-order delivery
  - ▼ Sequence numbers
- ▼ Congestion
  - ▼ Flow control, AIMD

## ■ Wait... What about UDP?


- ▼ Simple (de)multiplexing primitive over IP

# Layer encapsulation example

HTTP, Get  
index.html

**Application**


# Layer encapsulation example



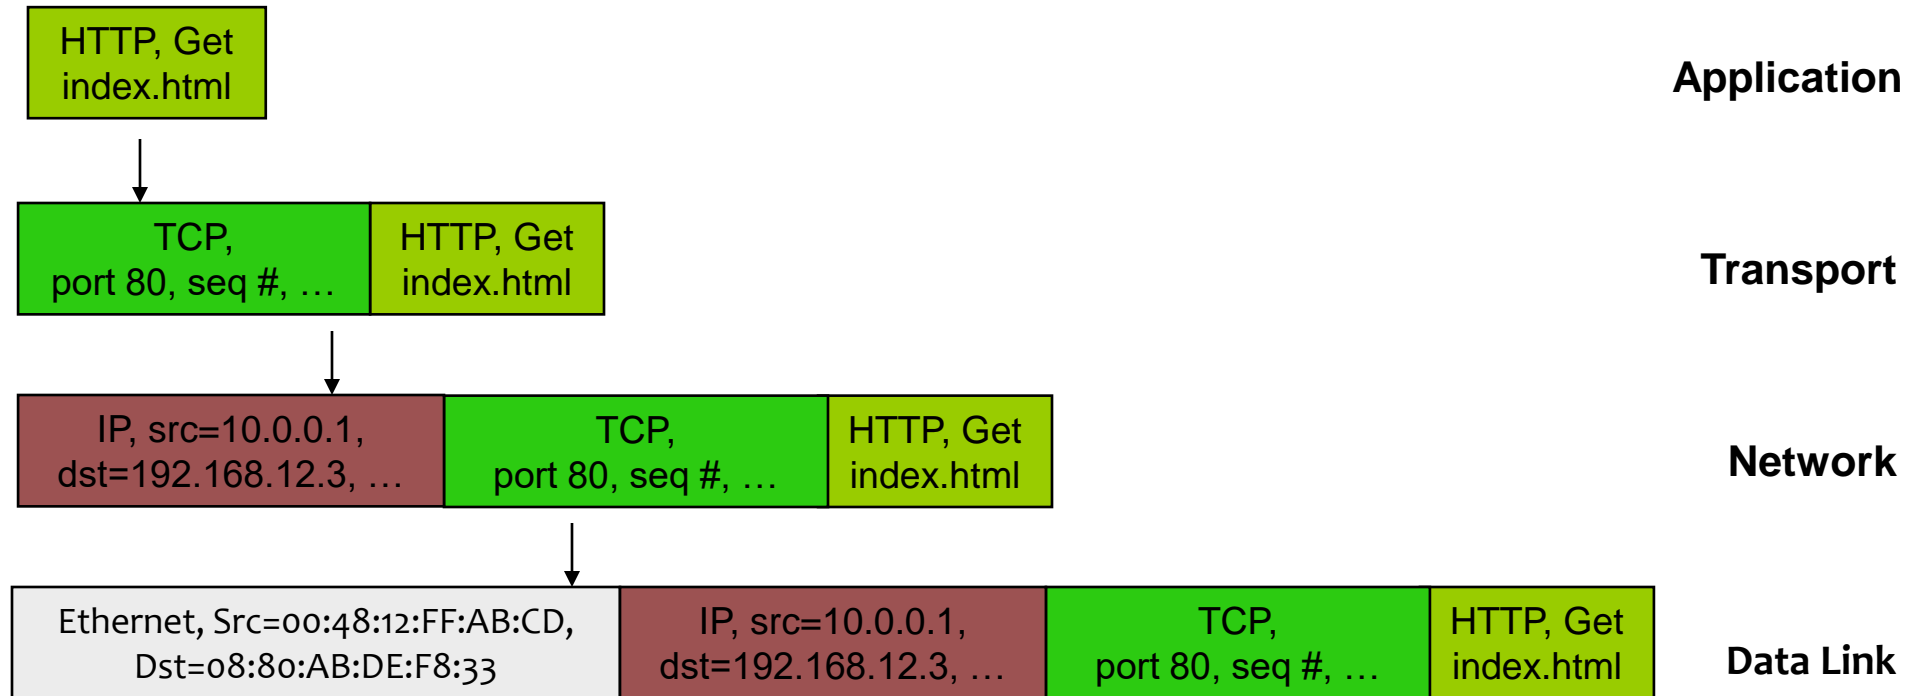
**Application**

**Transport**

# Layer encapsulation example

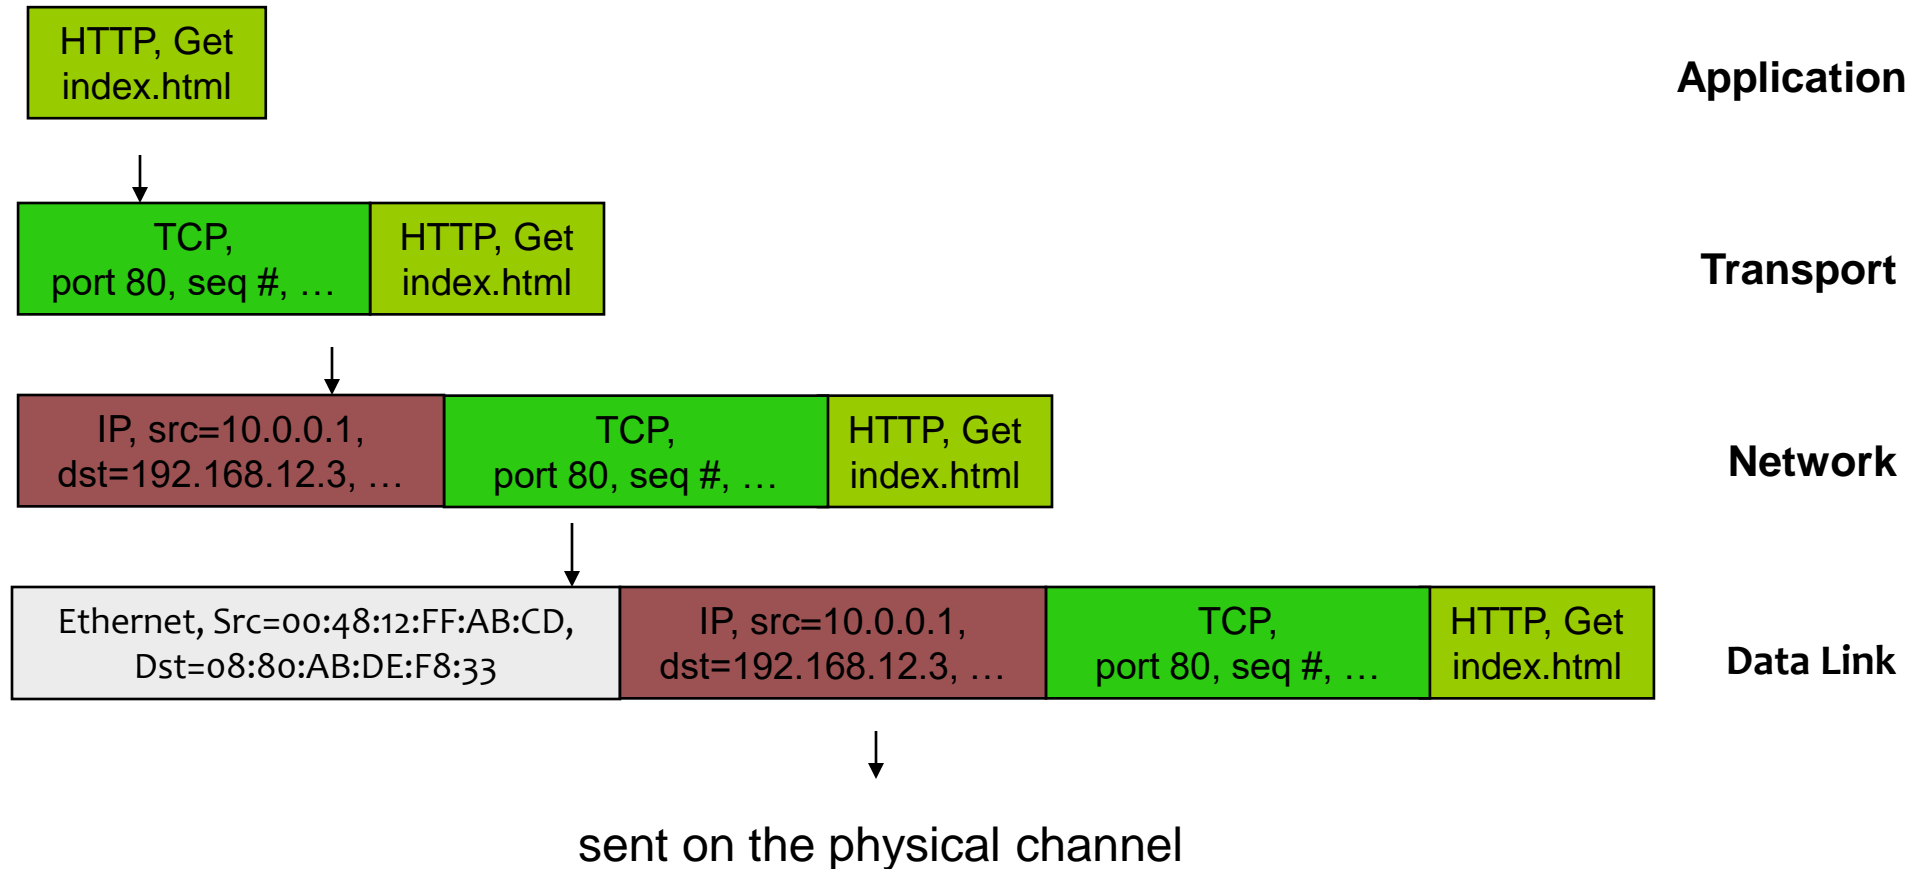


# Layer encapsulation example





# Layer encapsulation example



# Naming/Identification

## ■ IP addresses vs. readable host names

- ▼ Humans are not good at remembering numbers, use E.g., [www.cmu.edu](http://www.cmu.edu)
- ▼ Globally unique (but can correspond to multiple hosts, e.g., [www.google.com](http://www.google.com))


## ■ Naming system translates readable host names to physical address

- ▼ DNS translates name to IP address (e.g., 128.2.42.52)
- ▼ Address reflects location in the network

## ■ A Common authentication problem

- ▼ IP-based authentication/identification can be defeated by IP address spoofing

# Network security issues




- **Message authentication?**
- **How can S know that the packet really originated from A?**
- **Can B impersonate A to S?**
- **Can C impersonate A to S?**
- **Shared networks are easy to eavesdrop on**
  - ▼ E.g., just set interface in promiscuous mode  
tcpdump, ethereal
  - ▼ B can easily eavesdrop on the packet A sent to S
- **Note: eavesdropping or injection of arbitrary packets easy in today's networks**

# Security problems in the TCP/IP protocol suite

- By Steven M. Bellovin, 1989
- Classic network security paper
- Wakeup call for networking researchers
- Paper lists many security vulnerabilities
- Attacks studied in this lecture
  - ▼ IP-based authentication
  - ▼ TCP level attacks
  - ▼ Routing attacks
  - ▼ Application-level attacks
- More next time, and even more next semester


# Review: TCP 3-way handshake



# TCP initial sequence number (ISN) prediction attack

- **Objective of the attack: Impersonation of a communication participant**
  - ▼ E.g., Alice may talk to Bob
  - ▼ Mallory can pose as Alice, even though Mallory cannot intercept traffic between Alice and Bob
- **Risk: severe, particularly when authentication is based on IP address**
  - ▼ e.g., rsh: arbitrary execution of commands possible

# Abusing the TCP 3-way handshake



**(Mallory can send commands to Bob as Alice)**

# Predicting $ISN_{Bob}$


- **For the attack to be successful, Mallory needs to be able to predict  $ISN_{Bob}$** 
  - ▼ Easy on a LAN (just eavesdrop)
  - ▼ But not much harder on a WAN...
- **Flawed TCP ISN choices**
  - ▼ Always start at same ISN
  - ▼ After each connection,  $ISN++$
  - ▼  $ISN = (c1 + c2 * (\text{time in ms})) \bmod 2^{32}$




# Better ISN choices

- **ISN is should be unpredictable!**
- **ISN = rand() function of C library?**
  - ▼ How random is rand()?
  - ▼ Especially at boot time?
- **current ISN = H(previous ISN)?**
- **ISN = AES<sub>K</sub>(counter++)?**

# Doesn't Alice reset the connection?

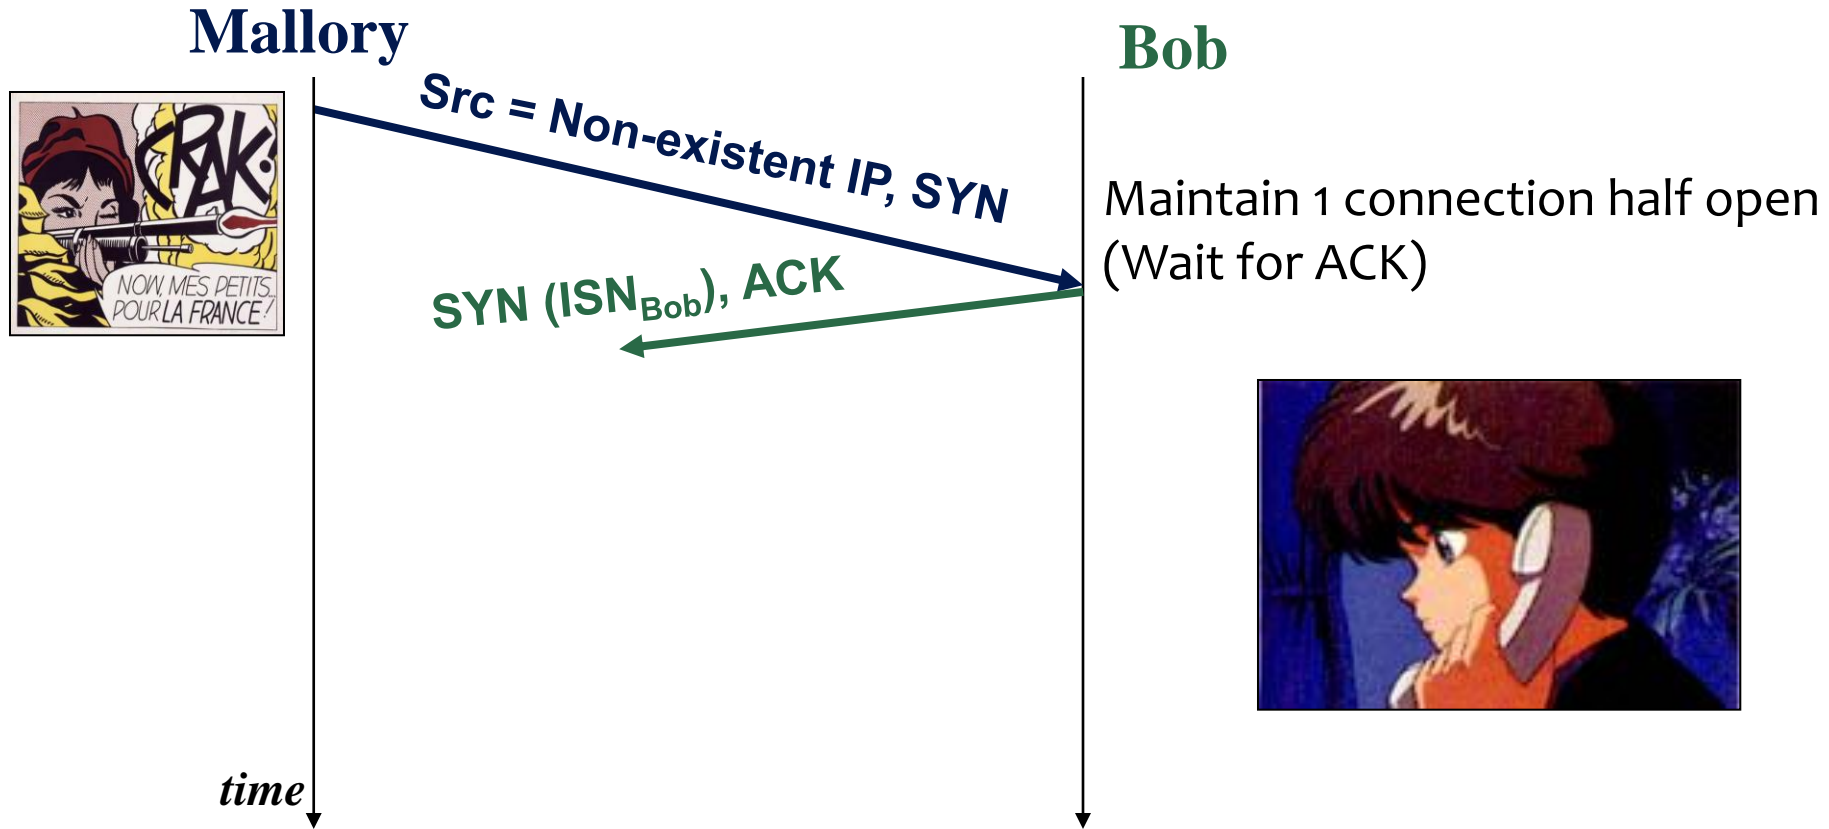


# Abusing the TCP 3-way handshake (revised)

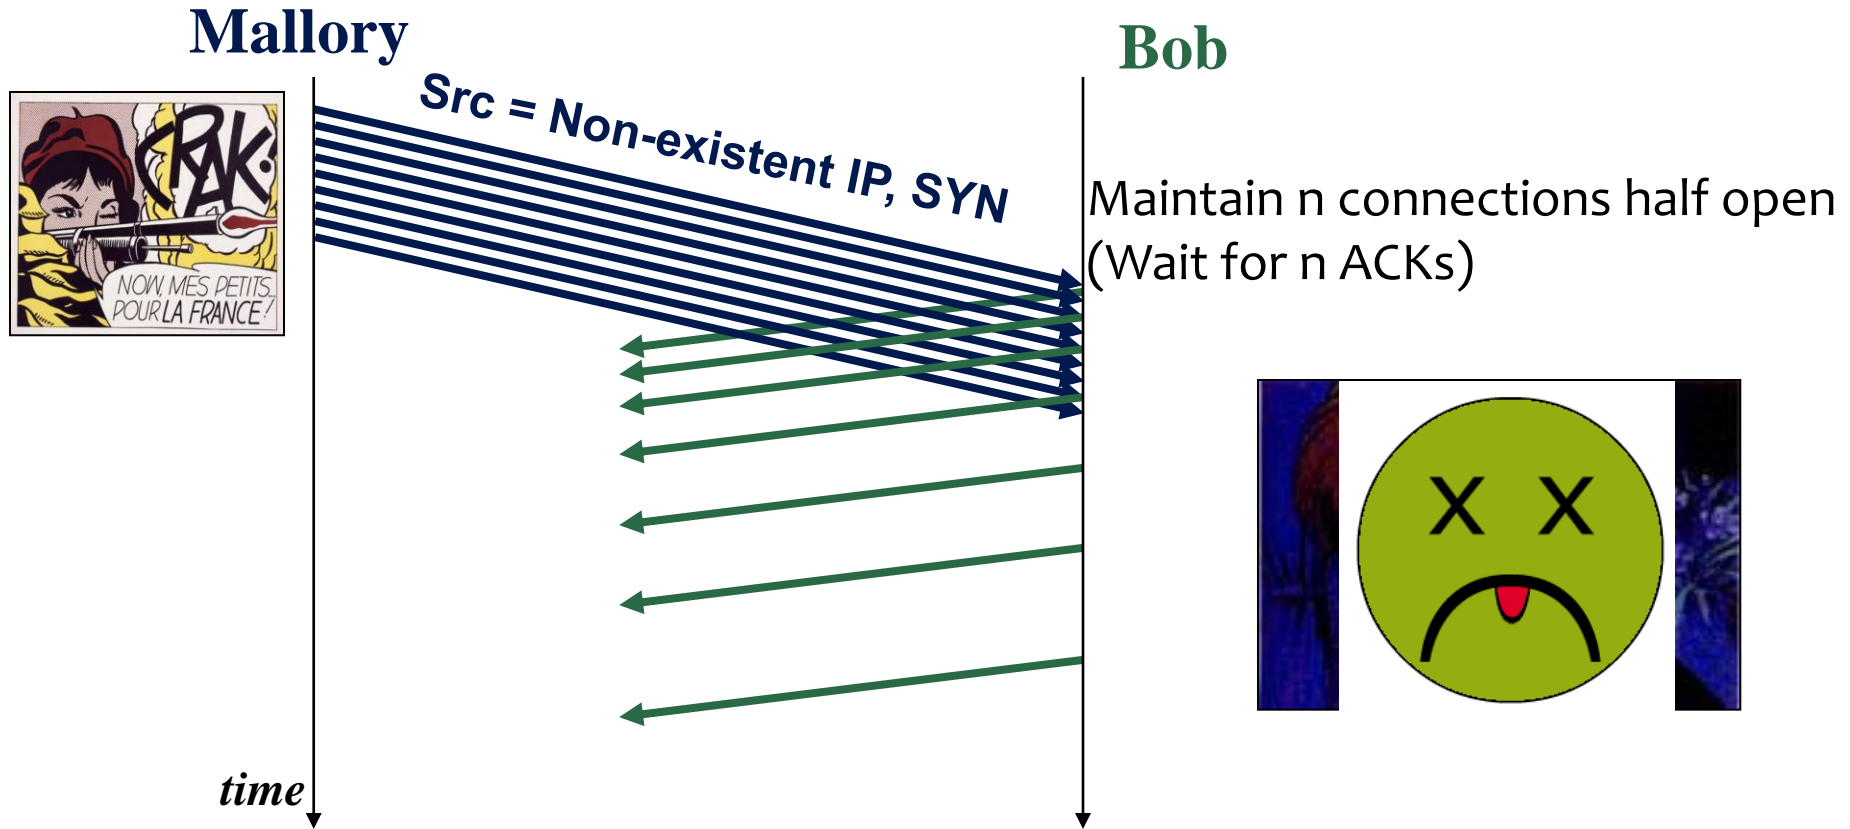


(Mallory can send commands to Bob as Alice)

# TCP SYN flood



# TCP SYN flood



# SYN flood details

## ■ Why does server exhaust resources?

- ▼ Need to store requests for 511 seconds
- ▼ Server has finite-size queue for incomplete connections, usually 1024 entries

## ■ Memory is cheap, why not store all requests?

- ▼ 736 bytes per connection...

## ■ Why store any information at all?

- ▼ If SYN ACK dropped by network, server re-sends SYN ACK until timeout or client sends ACK, otherwise legitimate clients will wait
- ▼ TCP options (performance enhancements) need to be stored, otherwise slow connection

# Defense against SYN flood: TCP SYN cookie

## ■ Approach

- ▼ Server computes server ISN (SYN cookie) based on connection state, reply to client, server keeps no state after SYN
- ▼ Client sends ACK with cookie
- ▼ Server verifies cookie, if correct, allocates connection state, reconstruct the connection state (with some information loss such as special TCP options)

## ■ What about: $\text{Cookie} = H(\text{SIP}, \text{CIP}, \text{Sport}, \text{Cport})$ ?

## ■ Better: $\text{Cookie} = H(\text{SIP}, \text{CIP}, \text{Sport}, \text{Cport}, \text{skey})$

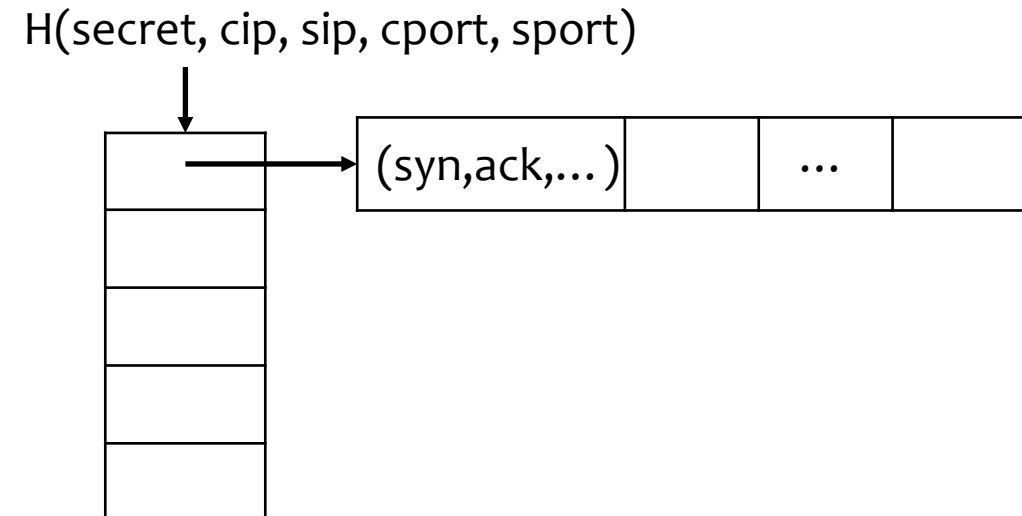
# Question

- **Connection state is much larger than incomplete (half-open) connection state**
- **Why does client not follow up with ACK packet in TCP SYN flooding?**
  - ▼ No IP source address spoofing
    - ▼ Server can easily identify attacker, since it has thousands of open connections
    - ▼ Server can block all connections from attacker IP
  - ▼ IP source address spoofing
    - ▼ Attacker does not receive SYN ACK, does not know server ISN, cannot send ACK



# TCP SYN cache

- **Best strategy: keep efficient SYN cache for incomplete connections, only if too many requests send SYN cookies**
- **SYN cache design**
  - ▼ Efficient hash table
  - ▼ Each hash table entry has linked list with all incomplete connections
  - ▼ Limit number of entries on linked list per bucket
  - ▼ Compute hash with a secret



# Results




Figure 3: Time needed to connect() to remote system.

# Tampering routing tables

- Routing tables updated from information received from topological neighbors
- Checks are non-existent or casual (IP address-based authentication)
- Possibility to route all traffic through a malicious host

# Defending against routing attacks

- **Authentication of all packets**
  - ▼ is difficult for broadcast protocols (e.g., RIP)
- **Do not accept new routes to your own networks**

# Application problems

- **Finger service gives out information useful for password cracking**
- **User name and password are often sent in the clear**
  - ▼ POP email retrieval protocol
  - ▼ FTP
  - ▼ Telnet
- **TFTP causes many weaknesses**
  - ▼ e.g., home routers, remote booting
- **DNS insecurity**

# Defenses

- **Authenticated communication prevents packet injection attacks**
- **Encrypted communication protects clear text**
- **IPsec, SSH, TLS/SSL provide secrecy and authentication**

# 1989... That's old, isn't it?

## ■ Aren't most of the attacks listed in the paper defeated by now?

### ▼ You wish...

- ▼ A number of network appliances (cheap, or less cheap) still have predictable TCP ISNs
- ▼ Some servers still run unencrypted POP/SMTP/IMAP or services like fingerd, r-services (rarer, but still)

### ▼ Even for attacks that have been addressed, we need to learn from our mistakes

- ▼ Source routing is evil, we disable it
- ▼ Great, but then we advocate MPLS, which may use source routing, overlay source routing, ...

# Take away slide

- **TCP/IP remarkable at fault-tolerance...**
- **... but not built with malicious attackers in mind**
  - ▼ To their credit, the designers of TCP/IP **did** foresee most of these problems
  - ▼ But crypto-based solutions couldn't easily make it into open standards back then
- **Be careful with what you assume from a protocol**
  - ▼ TCP and IP **do not provide any authentication**
    - ▼ An IP address is just that: an address, **not** an identity
  - ▼ Host-based authentication is an heresy
  - ▼ Need additional mechanisms (Kerberos, ...)
- **Relatively easy to abuse application-layer services**
  - ▼ Information services (SNMP, ...) are covert channels leaking information